

APL 405 Assignment-3

Nihal Pushkar

1.2 Mean of the Normalized 'MaxTemp'

```
✓ [186] X,y,mean = lr().data_clean(df_train)
0s      mean
      0.4647423708243247
```


1.4 Plotting cost as function of iteration

```
def minCostFun(self, w_ini, X_train, y_train, iters):
    # iters - Maximum no. of iterations; X_train - Numpy array
    lambda_ = 0.1 # Regularization parameter
    X_train = np.vstack([np.ones(len(X_train)),X_train.transpose()]).transpose() # Add '1' for bias term

    def f(w):
        return costing().costFunctionReg(w,X_train,y_train,iters)[0]

    li = []
    i = 1
    def callbackF(xi):
        global i
        li.append(f(xi))


    res = optimize.minimize(f,w_ini, method = 'TNC',callback = callbackF,options = {'maxiter' : iters})
    w_opt = res.x # Optimized weights rounded off to 3 decimal places
    p = costing().predict(w_opt,X_train)
    ones = np.ones(len(p))
    e = ones.transpose()@abs(p-y_train)/len(X_train)
    acrcy = (1-e)*100 # Training set accuracy (in %) rounded off to 3 decimal places
    return w_opt, acrcy,li
```



2.0 Load Data

```
# Load Data
datan1 = np.loadtxt('/content/drive/MyDrive/GitHub/APL405-1/Week_03_Logistic_Regression/nonLinearClass.txt', delimiter = ',')
x1 = datan1[:,1]
x0 = datan1[:,0]
y = datan1[:,2]
```

2.1 Visualization of the Data



2.2 Feature mapping

2.2 Feature mapping

From the plot, it is clear that the dataset can not be separated into positive and negative example by using a straight line through the plot.

- One way to fit the data better is to create more features from each data point. We can define a function `mapFeature` in order to map the features into all polynomial terms of x_1 and x_2 up to the n^{th} power.

$$\text{mapFeature}(x) = [1 \quad x_1 \quad x_2 \quad x_1^2 \quad x_1x_2 \quad x_2^2 \quad x_1^3 \quad \dots \quad x_1x_2^{n-1} \quad x_2^n]^T$$

- As a result of this mapping, the **vector of two features (the scores on two QA tests)** will be transformed into a multi-dimensional vector.
- A logistic regression classifier trained on this higher-dimension feature vector will have a more complex decision boundary and will appear nonlinear when drawn in our 2-dimensional plot.
- While the feature mapping allows us to build a more expressive classifier, it also more susceptible to overfitting.

```
def mapFeature(X1, X2, degree):
    out = []
    for i in range(degree+1):
        for j in range(i+1):
            t = (X1**j)*(X2**(i-j))
            out.append(t)
    return np.array(out)
```

```
# Note that mapFeature should also add a column of ones, so the intercept/bias term is handled
degree = 6
m = len(x0)
X = []
for i in range(m):
    X.append(mapFeature(x0[i], x1[i], degree))
X = np.array(X)
print(X.shape)

(118, 28)
```

2.3 Cost function and Gradient

2.3 Cost function and gradient

Now the code to compute the cost function and gradient for regularized logistic regression will be implemented in the function `costFunctionReg` below to return the cost and gradient.

Recalling that the regularized cost function in logistic regression is

$$J(w) = \frac{1}{m} \sum_{i=1}^m \left[-y^{(i)} \log(h_w(x^{(i)})) - (1 - y^{(i)}) \log(1 - h_w(x^{(i)})) \right] + \frac{\lambda}{2m} \sum_{j=1}^n w_j^2$$

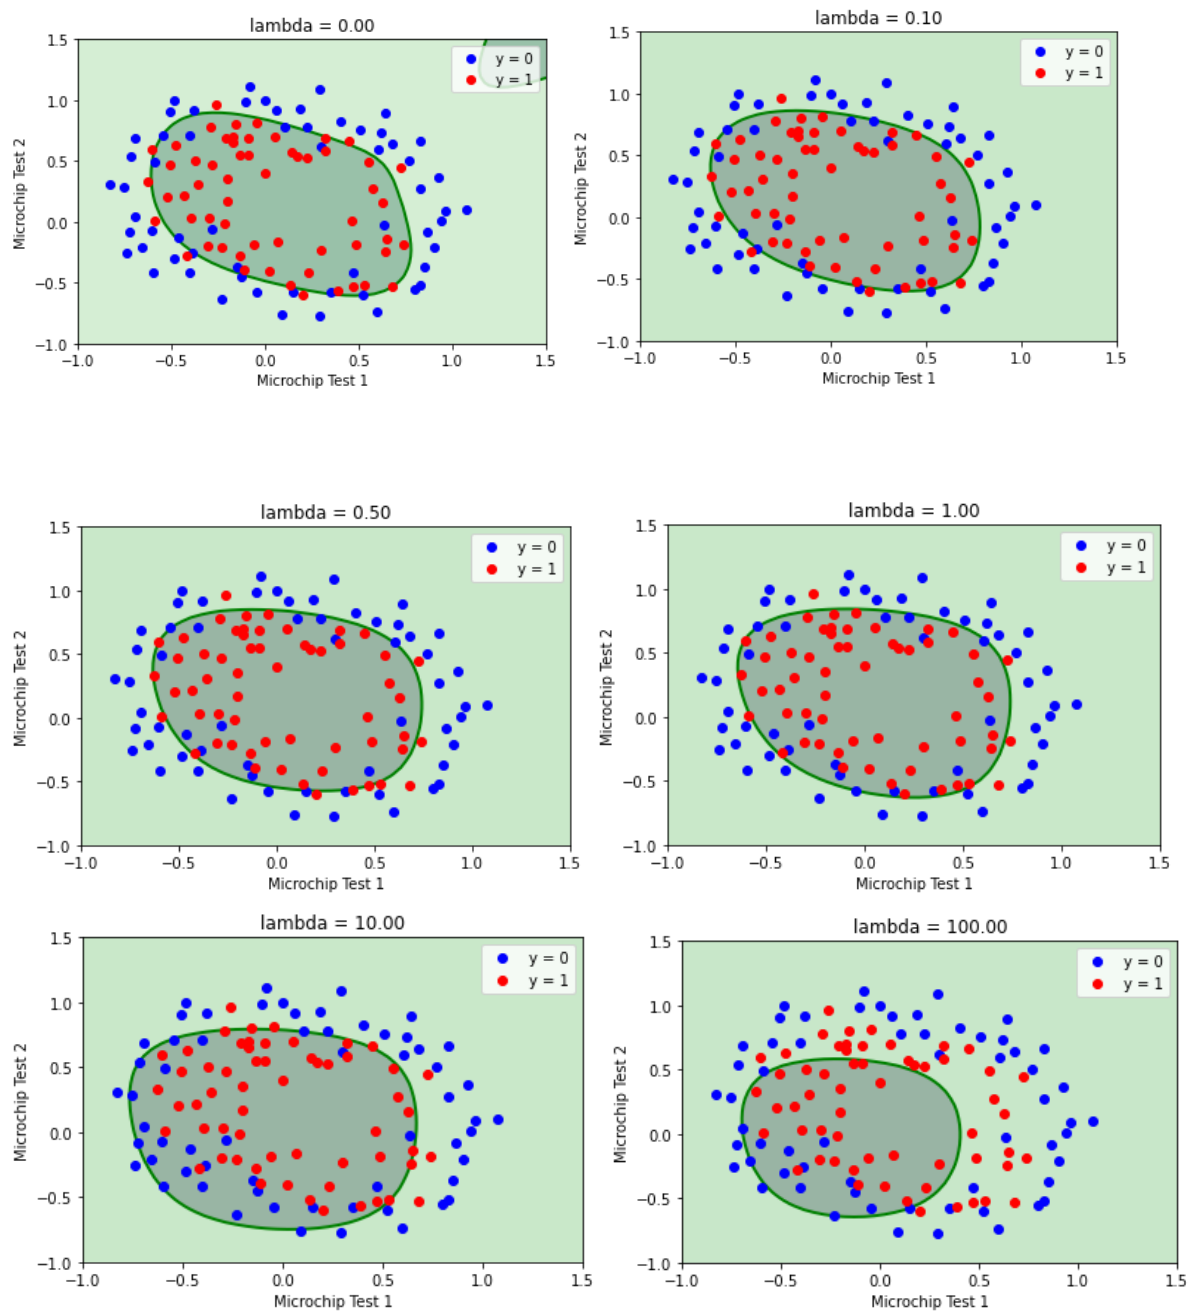
Note that one should not regularize the parameters w_0 . The gradient of the cost function is a vector where the j^{th} element is defined as follows:

$$\frac{\partial J(w)}{\partial w_0} = \frac{1}{m} \sum_{i=1}^m (h_w(x^{(i)}) - y^{(i)}) x_j^{(i)} \quad \text{for } j = 0$$
$$\frac{\partial J(w)}{\partial w_j} = \left(\frac{1}{m} \sum_{i=1}^m (h_w(x^{(i)}) - y^{(i)}) x_j^{(i)} \right) + \frac{\lambda}{m} w_j \quad \text{for } j \geq 1$$

```
[ ] def sigmoid(z):  
    g = 1/(1+np.exp(-1*z))  
    return g
```

```
def costFunctionReg(w, X, y, lambda_):  
    # =====  
    p = len(w)    # no. of parameters, order(X) = m*p  
    m = len(X)    # no. of training examples  
    h = sigmoid(X@w)    # (m*p)*(p*1) = (m*1)  
    sum1 = y.transpose()@np.log(h) + (1-y).transpose()@np.log(1-h)  
    w2 = w.copy()  
    w2[0] = 0  
    sum2 = (lambda_/(2*m))* (w2.transpose()@w2)  
    J = (-1*sum1/m) + sum2    # Cost 'J' should be a scalar  
    grad = (1/m)*((X.transpose())@(h-y) + lambda_*w2)  
    # =====  
    return J, grad
```

2.4 Plots with variation of the regularization term: $\lambda = 0, 0.1, 0.5, 1, 10, 100$



For $\lambda = 0$ we observe that the model tries to get every positive training set. But while doing so the curves become sharp. Case of overfitting

Whereas for $\lambda = 0.1$ we observe the curve to be smooth enough and being covering positive training points.

For $\lambda = 0.5, 1$ curve is quite smooth but misses more points as compared to the case of $\lambda = 0.1$, but still defines the decision boundary good enough.

Whereas for $\lambda = 10, 100$ the weights were so suppressed to reduce the cost that it couldn't even define the decision boundary properly.